

05/13

Concurrent Programming

Concurrent Programming is hard → 왜 어렵나?

Races

→ 결과가 임의의 scheduling 결정에 의해 좌우됨.

Deadlock

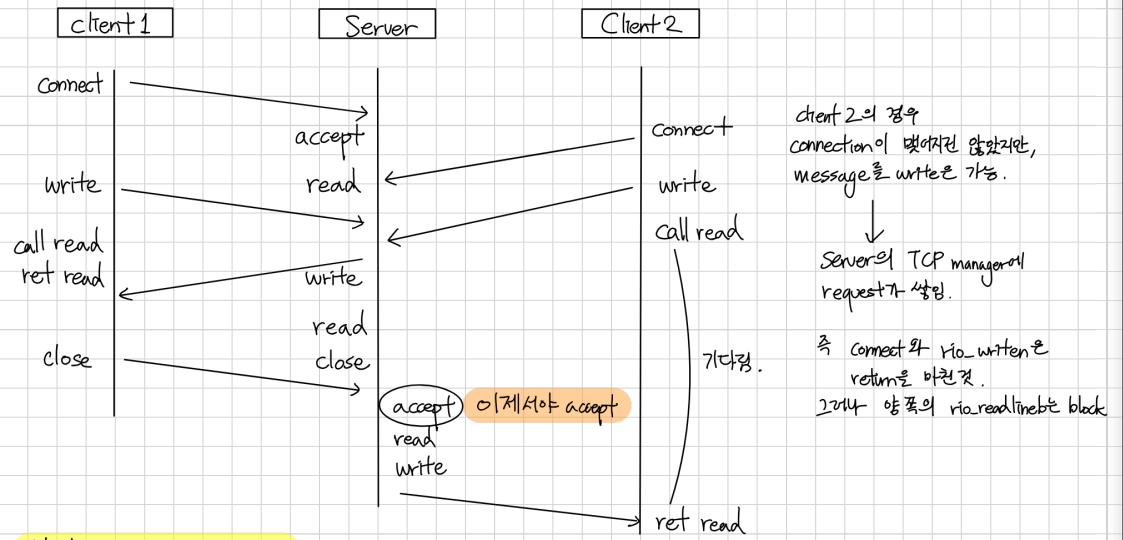
→ 부정적인 resource allocation이 프로세스의 진행 막음.

→ 서로 다른 프로세스가 서로의 진행을 기다림. ex) signal handler의 printf.

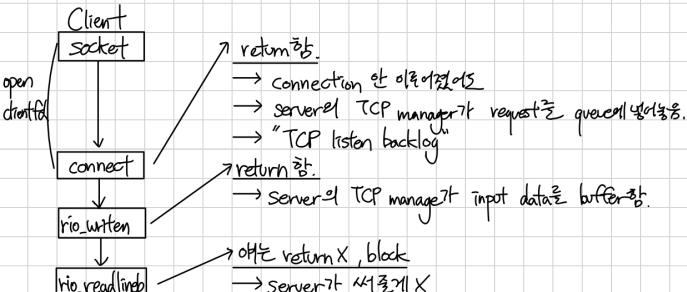
Livelock / Starvation / Fairness

→ 반복적 유품이나가 실패? 상태는 변화하는데 진전이 X

Problem of Iterative Servers (Not Concurrent Servers)



Why Does Second Client Block?



Approaches for Concurrent Servers

1) Process-based

- 운영체제 fork를 이용, 여러 별개 프로세스 사용
- Kernel이 automatically multiple logical flows 운영
- 각 flows는 각각의 private address space

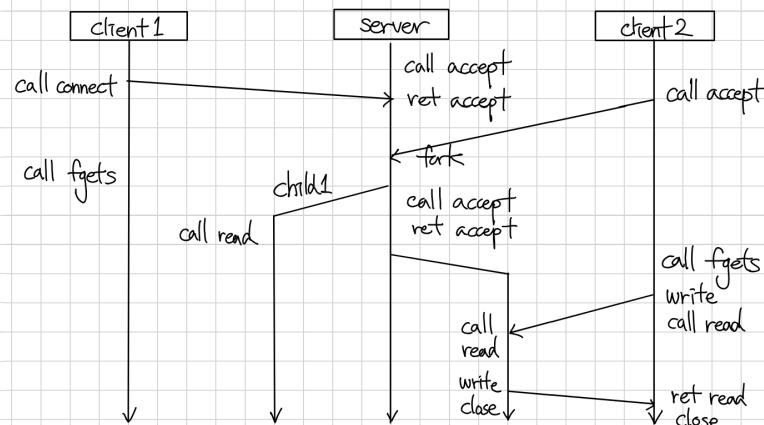
2) Event-Based

- 운영체제에서 일련의 이벤트 처리 흐름 이용
- Programmer가 manually multiple logical flows 운영
- 모든 flows가 same address space 공유
- I/O multiplexing 활용

3) Thread-based

- 운영체제에서 thread 만들
- Kernel이 automatically multiple logical flows 운영
- 각 flow가 same address space 공유
- process based와 event based의 혼합

Approach #1: Process-based Servers

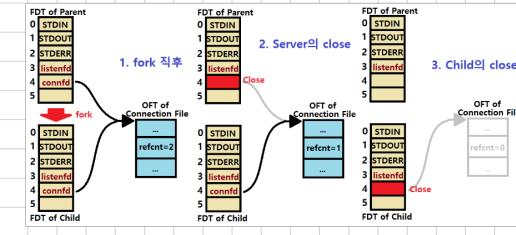


```

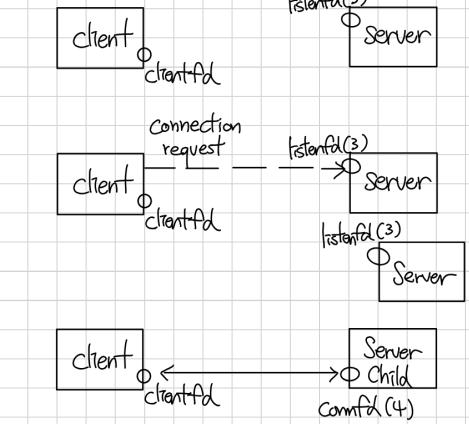
int main(int argc, char **argv)
{
    int listenfd, connfd;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;

    Signal(SIGCHLD, sigchld_handler);
    listenfd = Open_listenfd(argv[1]);
    while (1) {
        clientlen = sizeof(struct sockaddr_storage);
        connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
        if (Fork() == 0) {
            Close(listenfd); /* Child closes its listening socket */
            echo(connfd); /* Child services client */
            Close(connfd); /* Child closes connection with client */
            exit(0); /* Child exits */
        }
        Close(connfd); /* Parent closes connected socket (important!) */
    }
}
void sigchld_handler(int sig)
{
    while (waitpid(-1, 0, WNOHANG) > 0)
        ;
    return;
}

```



Concurrent Server: accept illustrated

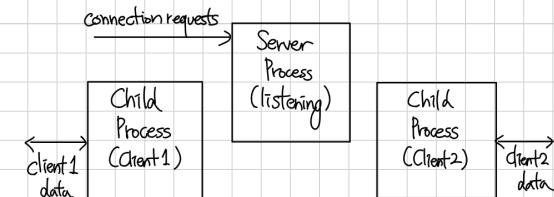


1. Server는 accept()가 block된 상태.
listenfd의 request가 읽기까지 wait

2. client가 connect 함수를 호출하여
connection request를 보냄.

3. server는 accept를 통해 connfd를 return, connection 완성!!

Process-based Server Execution Model



- 각 client는 child process가 처리.
- 각 child끼리 어떤 상태, 공간도 공유X
- Parent는 connfd를 close
Child는 listenfd를 close

Issues with Process-based Servers

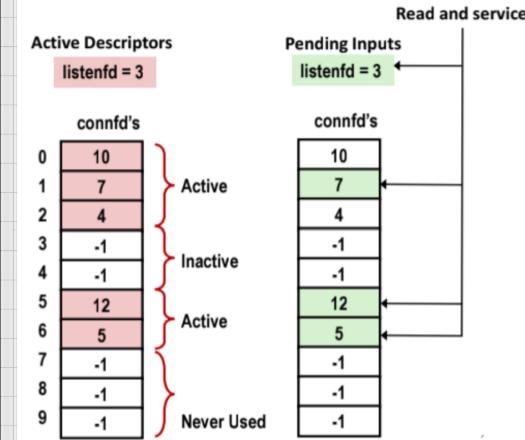
- Listening Server process는 zombie children을 반드시 reaping 해야함.
- Parent process는 모든 connfd를 shutdown.
fork 이후 refcnt(connfd)=2
connection은 refcnt(connfd)=0이 되기 전까지 닫히지X

Pros and Cons of Event-based Servers

- ⊕ handle multiple connections concurrently
- ⊕ clean sharing model
 - descriptors (no, 공유) 왜? 바로 close
 - file tables (yes)
 - global variables (no)
- ⊕ simple & straightforward
- ⊖ additional overhead for process control
- ⊖ nontrivial to share data between processes
 - requires IPC

Approach #2 Event-based Concurrent Server

- Server가 'Active Connection'에 대한 집합 만들.
- ↳ 'Array of connfd's'
- 어떤 descriptor pending input이 있는지 조사.
- 만약 listenfd의 pending input → Connection accept 수행 (array of all connfd 초기화)
- Connfd all " " ⇒ 서비스 처리



I/O Multiplexing-using select

```
int select(int n, fd_set *readfd, fd_set *writefd, fd_set *exceptfd, struct timeval *timeout)
    ⇒ active(ready) descriptor의 개수를 반환, timeout
```

```
void FD_ZERO(fd_set *fdset)
void FD_CLR(int fd, fd_set *fdset)
void FD_SET(int fd, fd_set *fdset)
int FD_ISSET(int fd, fd_set *fdset)
```

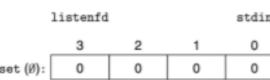
bit vector 0으로 초기화.
fd all 위치한 k번째 bit 0으로.
" k번째 bit setting.
k번째 bit가 set인지 확인

ex) stdin(0번)에 1을 set, listenfd(3번)에 1을 set하면, 0과 3이 pending input이 있는지 확인 가능.

```

1 #include "csapp.h"
2 void echo(int connfd);
3 void command(void);
4
5 int main(int argc, char **argv)
6 {
7     int listenfd, connfd;
8     socklen_t clientlen;
9     struct sockaddr_storage clientaddr;
10    fd_set read_set, ready_set; ※ 초기화
11
12    if (argc != 2) {
13        fprintf(stderr, "usage: %s <port>\n", argv[0]);
14        exit(0);
15    }
16    listenfd = Open_listenfd(argv[1]);
17
18    FD_ZERO(&read_set); /* Clear read set */ → 모든 false로 초기화
19    FD_SET(STDIN_FILENO, &read_set); /* Add stdin to read set */ → stdin 활성화
20    FD_SET(listenfd, &read_set); /* Add listenfd to read set */ → listenfd 활성화
21

```

line 18: read_set (0):


line 19-20: read_set ({0,3}):


커버드 대기 중 처리.

```

he ready set
line 24: listenfd      stdin
          3       2       1       0
ready_set ({0}): 0       0       0       1
→ stdin이나 listenfd가 대기 중인 이벤트가 발생한 후 ready_set은 {0}이 됨.

```

```

22    while (1) {
23        ready_set = read_set; ready_set = read_set; → ready_set은 원본 대기 목록을 계속 유지해야 함.
24        Select(listenfd+1, &ready_set, NULL, NULL, NULL);
25        if (FD_ISSET(STDIN_FILENO, &ready_set)) → stdin 대기 중인 이벤트
26            command(); /* Read command line from stdin */
27        if (FD_ISSET(listenfd, &ready_set)) { → listenfd 대기 중인 이벤트
28            clientlen = sizeof(struct sockaddr_storage);
29            connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
30            echo(connfd); /* Echo client input until EOF */
31            Close(connfd);
32        }
33    }
34
35
36    void command(void) {
37        char buf[MAXLINE];
38        if (!Fgets(buf, MAXLINE, stdin))
39            exit(0); /* EOF */
40        printf("%s", buf); /* Process the input command */
41    }

```

Issues with select.c

+ ready_set과 read_set 서로 다른 이유
 read_set은 원본 대기 목록을 계속 유지해야 함.
 그래서 select는 호출하면, fd_set을 변경 ⇒ 이벤트 발생한 FD만 고장기고, 나머지 0으로 만듬.

```

#include "csapp.h"

typedef struct { /* Represents a pool of connected descriptors */
    int maxfd;           /* Largest descriptor in read_set */
    fd_set read_set;    /* Set of all active descriptors */
    fd_set ready_set;   /* Subset of descriptors ready for reading */
    int nready;          /* Number of ready descriptors from select */
    int maxi;            /* High water index into client array */
    int clientfd[FD_SETSIZE]; /* Set of active descriptors */
    rio_t clientrio[FD_SETSIZE]; /* Set of active read buffers */
} pool;

int byte_cnt = 0; /* Counts total bytes received by server */

int main(int argc, char **argv)
{
    int listenfd, connfd;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
    static pool pool;

    if (argc != 2) {
        fprintf(stderr, "usage: %s <port>\n", argv[0]);
        exit(0);
    }
    listenfd = Open_listenfd(argv[1]);
    init_pool(listenfd, &pool);
}

while (1) {
    /* Wait for listening/connected descriptor(s) to become ready */
    pool.ready_set = pool.read_set;
    pool.nready = Select(pool.maxfd+1, &pool.ready_set, NULL, NULL, NULL);

    /* If listening descriptor ready, add new client to pool */
    if (FD_ISSET(listenfd, &pool.ready_set)) {
        clientlen = sizeof(struct sockaddr_storage);
        connfd = Accept(listenfd, (SA *)&clientaddr, &clientlen);
        add_client(connfd, &pool);
    }
    /* Echo a text line from each ready connected descriptor */
    check_clients(&pool);
}

```

```

1 void init_pool(int listenfd, pool *p)
2 {
3     /* Initially, there are no connected descriptors */
4     int i;
5     p->maxi = -1;
6     for (i=0; i< FD_SETSIZE; i++)
7         p->clientfd[i] = -1;
8
9     /* Initially, listenfd is only member of select read set */
10    p->maxfd = listenfd;
11    FD_ZERO(&p->read_set);
12    FD_SET(listenfd, &p->read_set);
13 }

```

Pool of connected descriptors

```

struct {
    int maxfd;
    fd_set read_set;
    fd_set ready_set;
    int nready;
    int maxi;
    int clientfd[FD_SETSIZE];
    rio_t clientrio[FD_SETSIZE];
} pool;

```

→ block이 안된다? → 요청이 온 상태인가?

```

1 void add_client(int connfd, pool *p)
2 {
3     int i;
4     p->nready--;
5     for (i = 0; i < FD_SETSIZE; i++) /* Find an available slot */
6         if (p->clientfd[i] < 0) {
7             /* Add connected descriptor to the pool */
8             p->clientfd[i] = connfd;
9             Rio_readinitb(&p->clientrio[i], connfd);
10
11            /* Add the descriptor to descriptor set */
12            FD_SET(connfd, &p->read_set);
13
14            /* Update max descriptor and pool high water mark */
15            if (connfd > p->maxfd)
16                p->maxfd = connfd;
17            if (i > p->maxi)
18                p->maxi = i;
19            break;
20        }
21    if (i == FD_SETSIZE) /* Couldn't find an empty slot */
22        app_error("add_client error: Too many clients");
23 }

1 void check_clients(pool *p)
2 {
3     int i, connfd, n;
4     char buf[MAXLINE];
5     rio_t rio;
6
7     for (i = 0; (i <= p->maxi) && (p->nready > 0); i++) {
8         connfd = p->clientfd[i];
9         rio = p->clientrio[i];
10
11        /* If the descriptor is ready, echo a text line from it */
12        if ((connfd > 0) && (FD_ISSET(connfd, &p->ready_set))) {
13            p->nready--;
14            if ((n = Rio_readlineb(&rio, buf, MAXLINE)) != 0) {
15                byte_cnt += n;
16                printf("Server received %d (%d total) bytes on fd %d\n",
17                      n, byte_cnt, connfd);
18                Rio_writen(connfd, buf, n);
19            }
20
21            /* EOF detected, remove descriptor from pool */
22            else {
23                Close(connfd);
24                FD_CLR(connfd, &p->read_set);
25                p->clientfd[i] = -1;
26            }
27        }
28    }
29}

```

code/conc/echoservers.c

Pros and Cons of Event-based Servers

- + One logical control flow and address space.
- + Can single-step with a debugger.
- + No process or thread control overhead.
 - Design of choice for high-performance Web servers and search engines. e.g., Node.js, nginx, Tornado
- - Significantly more complex to code than process- or thread-based designs.
- ~~- Hard to provide fine-grained concurrency~~
 - E.g., how to deal with partial HTTP request headers
- - Cannot take advantage of multi-core
 - Single thread of control